



# Introduction to Docker




---

# Agenda

What is Docker

Why should I care

How does it work



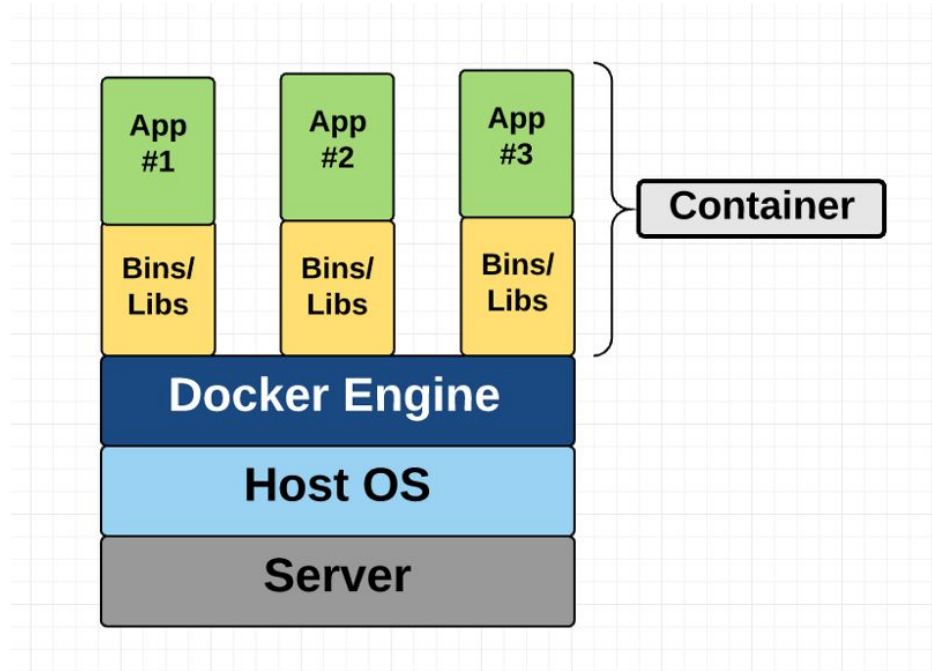

# What is Docker?

Docker is a platform for developing, shipping and running applications

- separates applications from the infrastructure
- manages infrastructure the same way you manage code
- uses "containers" to isolate application dependencies



# What is Docker?




---

# Agenda

What is Docker

Why should I care

How does it work



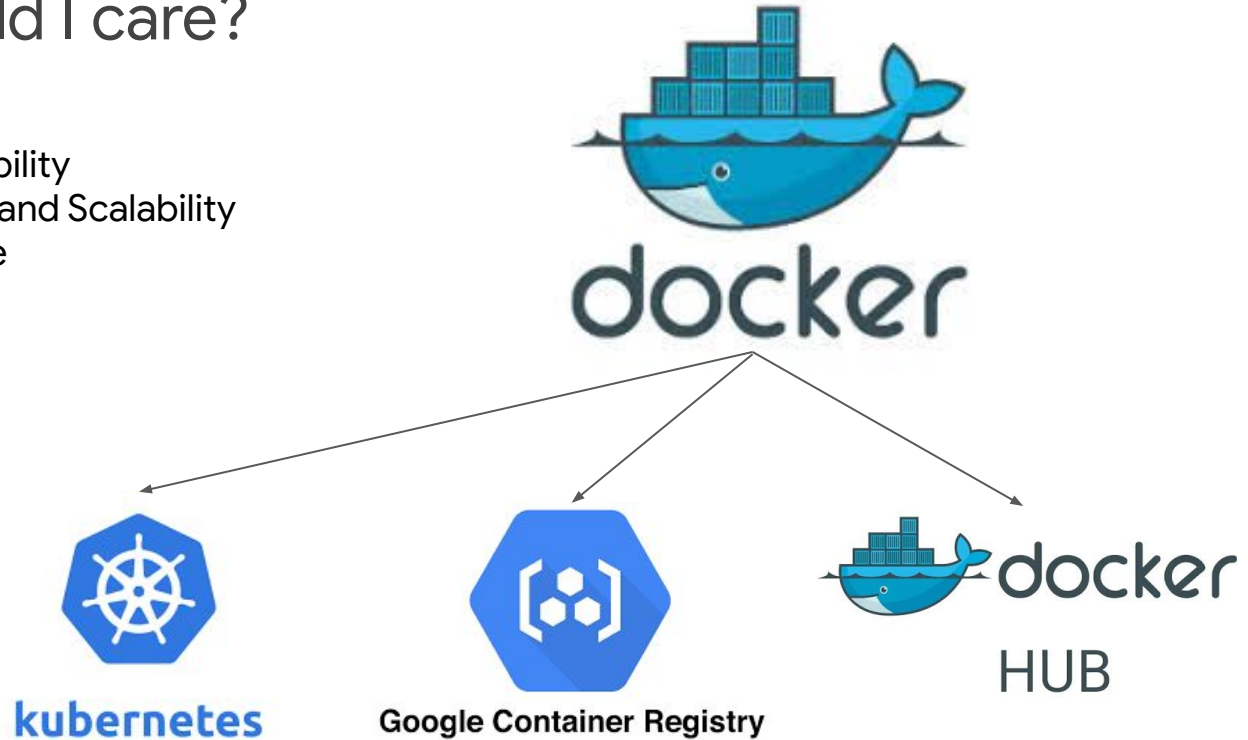
# Why should I care?

- Reproducibility
- Portability and Scalability



# Why should I care?

- Reproducibility
- Portability and Scalability
- Easy to use
- Speed




---

# Agenda

What is Docker

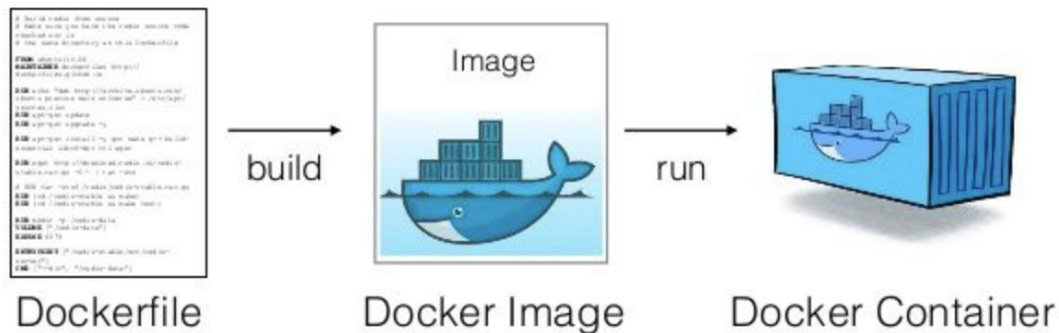
Why should I care

How does it work





# How does it work?



A Dockerfile is a recipe for creating an Image

A Docker Image is a serialized package containing an OS, the software dependencies, and the application code to be run.

A Docker Container is a process running an image

# How does it work?



Dockerfile

```
# Dockerfile
FROM python:2.7.15-jessie
WORKDIR /root

# Installs pytorch and torchvision.
RUN pip install torch==1.0.0 torchvision==0.2.1

# Installs cloudml-hypertune for hyperparameter tuning.
RUN pip install cloudml-hypertune

# Path configuration
ENV PATH $PATH:/root/tools/google-cloud-sdk/bin

# Make sure gsutil will use the default service account
RUN echo '[GoogleCompute]\nservice_account = default' >
/etc/boto.cfg

# Copies the trainer code
RUN mkdir /root/trainer
COPY trainer/mnist.py /root/trainer/mnist.py

# Sets up the entry point to invoke the trainer.
ENTRYPOINT ["python", "trainer/mnist.py"]
```

# How does it work?

base image you want to build  
on top of

# Dockerfile

FROM python:2.7.15-jessie

WORKDIR /root

# Installs pytorch and torchvision.

RUN pip install torch==1.0.0 torchvision==0.2.1

# Installs cloudml-hypertune for hyperparameter tuning.

RUN pip install cloudml-hypertune

# Path configuration

ENV PATH \$PATH:/root/tools/google-cloud-sdk/bin

# Make sure gsutil will use the default service account

RUN echo '[GoogleCompute]\nservice\_account = default' >  
/etc/boto.cfg

# Copies the trainer code

RUN mkdir /root/trainer

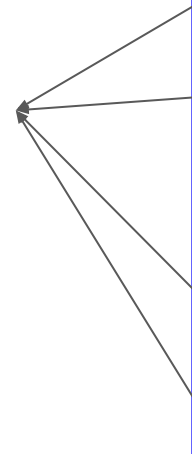
COPY trainer/mnist.py /root/trainer/mnist.py

# Sets up the entry point to invoke the trainer.

ENTRYPOINT ["python", "trainer/mnist.py"]

# How does it work?

run arbitrary shell commands  
to install packages and  
dependencies



```
# Dockerfile
FROM python:2.7.15-jessie
WORKDIR /root

# Installs pytorch and torchvision.
RUN pip install torch==1.0.0 torchvision==0.2.1

# Installs cloudml-hypertune for hyperparameter tuning.
RUN pip install cloudml-hypertune

# Path configuration
ENV PATH $PATH:/root/tools/google-cloud-sdk/bin

# Make sure gsutil will use the default service account
RUN echo '[GoogleCompute]\nservice_account = default' > /etc/boto.cfg

# Copies the trainer code
RUN mkdir /root/trainer
COPY trainer/mnist.py /root/trainer/mnist.py

# Sets up the entry point to invoke the trainer.
ENTRYPOINT ["python", "trainer/mnist.py"]
```

# How does it work?

change environments

```
# Dockerfile
FROM python:2.7.15-jessie
WORKDIR /root

# Installs pytorch and torchvision.
RUN pip install torch==1.0.0 torchvision==0.2.1

# Installs cloudml-hypertune for hyperparameter tuning.
RUN pip install cloudml-hypertune

# Path configuration
ENV PATH $PATH:/root/tools/google-cloud-sdk/bin


# Make sure gsutil will use the default service account
RUN echo '[GoogleCompute]\nservice_account = default' >
/etc/boto.cfg

# Copies the trainer code
RUN mkdir /root/trainer
COPY trainer/mnist.py /root/trainer/mnist.py

# Sets up the entry point to invoke the trainer.
ENTRYPOINT ["python", "trainer/mnist.py"]
```

# How does it work?

copy code from training  
application to container



```
# Dockerfile
FROM python:2.7.15-jessie
WORKDIR /root

# Installs pytorch and torchvision.
RUN pip install torch==1.0.0 torchvision==0.2.1

# Installs cloudml-hypertune for hyperparameter tuning.
RUN pip install cloudml-hypertune

# Path configuration
ENV PATH $PATH:/root/tools/google-cloud-sdk/bin

# Make sure gsutil will use the default service account
RUN echo '[GoogleCompute]\nservice_account = default' >
/etc/boto.cfg

# Copies the trainer code
RUN mkdir /root/trainer
COPY trainer/mnist.py /root/trainer/mnist.py

# Sets up the entry point to invoke the trainer.
ENTRYPOINT ["python", "trainer/mnist.py"]
```

# How does it work?

```
# Dockerfile
FROM python:2.7.15-jessie
WORKDIR /root

# Installs pytorch and torchvision.
RUN pip install torch==1.0.0 torchvision==0.2.1

# Installs cloudml-hypertune for hyperparameter tuning.
RUN pip install cloudml-hypertune

# Path configuration
ENV PATH $PATH:/root/tools/google-cloud-sdk/bin

# Make sure gsutil will use the default service account
RUN echo '[GoogleCompute]\nservice_account = default' > /etc/boto.cfg

# Copies the trainer code
RUN mkdir /root/trainer
COPY trainer/mnist.py /root/trainer/mnist.py

# Sets up the entry point to invoke the trainer.
ENTRYPOINT ["python", "trainer/mnist.py"]
```

configures entrypoint



# Building, Pushing, and Running a Docker image

```
IMAGE_NAME=gcr.io/${PROJECT_ID}/node-app:0.2
```

```
docker build -t $IMAGE_NAME .
```

Supposes the Dockerfile is in the CWD (".")



```
docker push $IMAGE_NAME
```

```
docker run -p 4000:80 $IMAGE_NAME
```

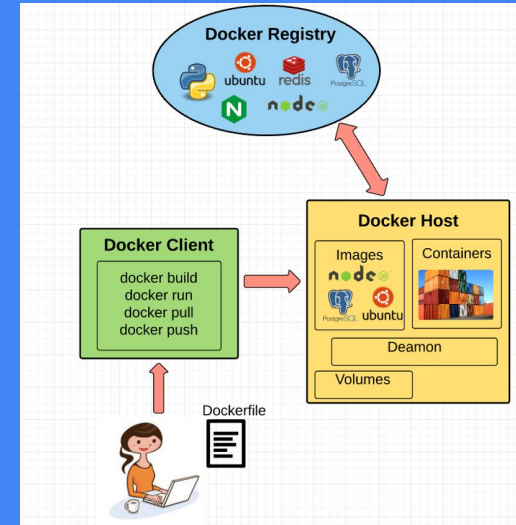


# Lab

In this lab, you will learn how to:

- How to build, run, and debug Docker containers.
- How to pull Docker images from Docker Hub and Google Container Registry.
- How to push Docker images to Google Container

[docker\\_and\\_kubernetes/labs/1\\_intro\\_docker.ipynb](https://docker_and_kubernetes/labs/1_intro_docker.ipynb)



cloud.google.com

